Introducing watsonx.governance

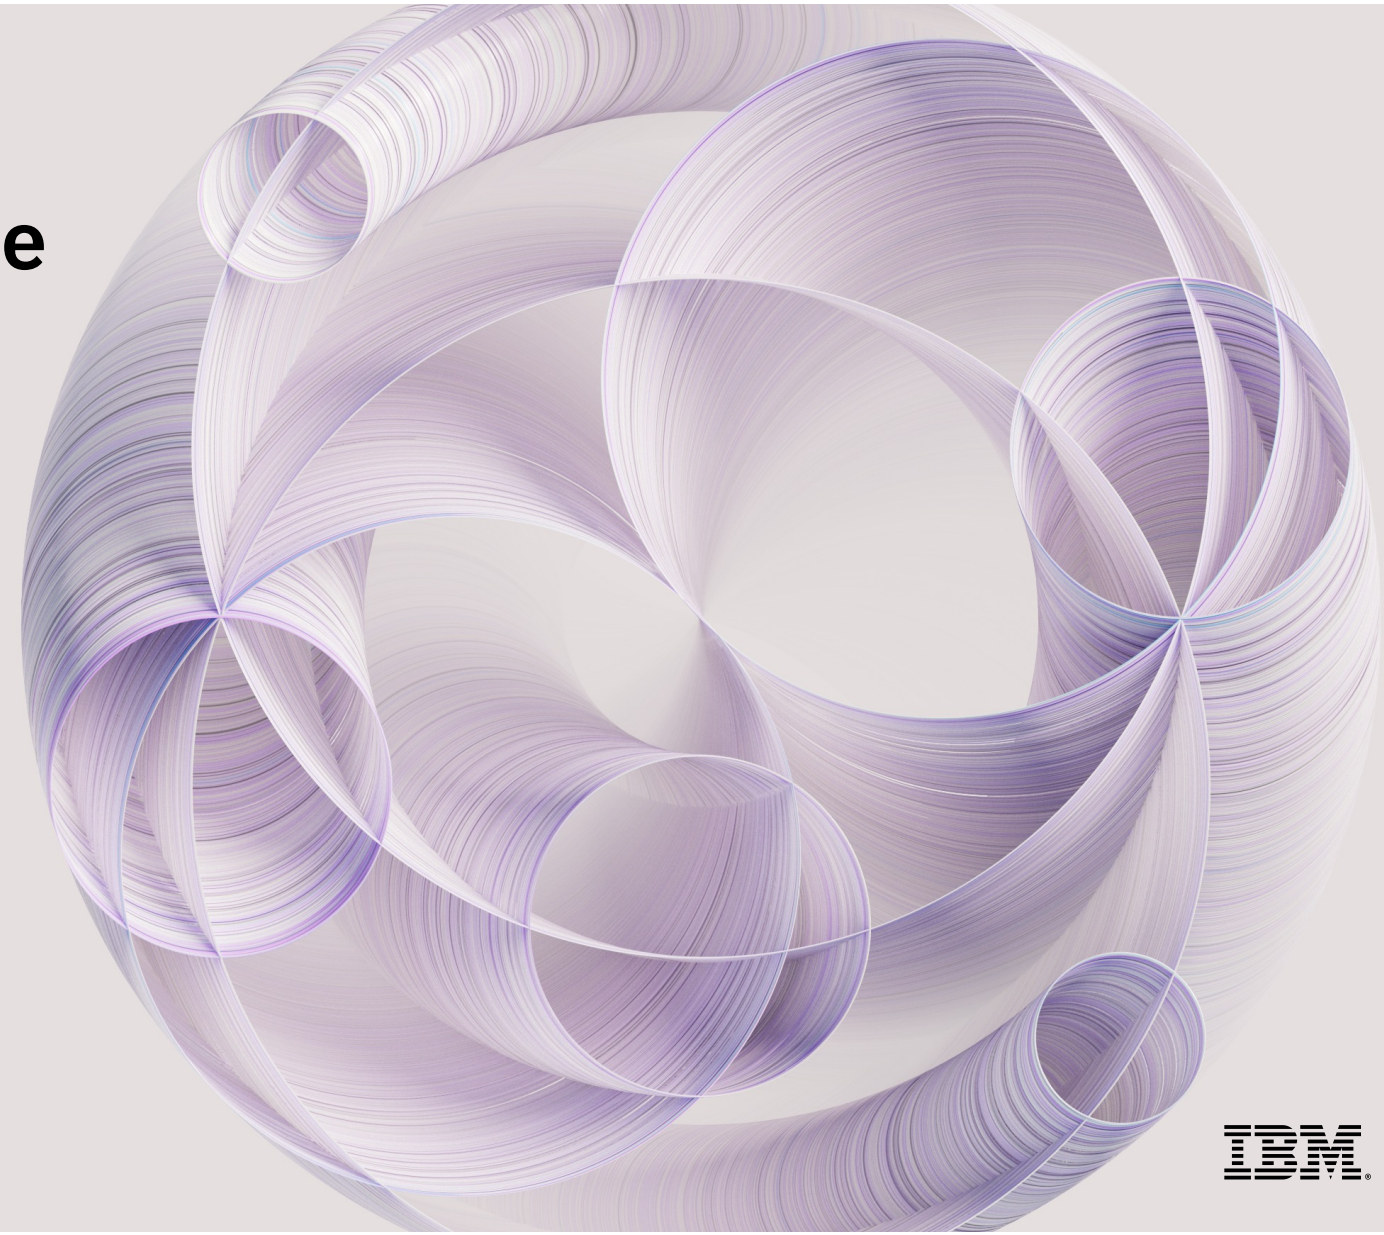

The promise of AI is clear, but implementations come with questions

IDC predicts AI lifecycle opportunity is 6.7B in 2022 with AI Governance opportunity growing at 55%

How do we operationalize AI with confidence?

"**Protected/sensitive data like ethnicity and gender** should not be used by AI models for hiring. How can we ensure fairness ?"

How do I enable responsible use of AI to manage risk?


RETAIL OCTOBER 10, 2018 / 4:04 PM / UPDATED 2 YEARS AGO

Amazon scraps secret AI recruiting tool that showed bias against women

How do I protect and scale against regulations?

Reputation: poor governance can damage consumer trust

Organizations need to:

- Only use personal information when it is needed and with the user's consent
- Ensure results are free from bias (ex. sex, age, etc) and guided by ethical and transparent principles
- Build consumer confidence, providing explanations for business decisions (ex. credit, membership denials)
- Proactively detect fraud and risk to consumer’s accounts

Risk: bad press leading to loss of revenues

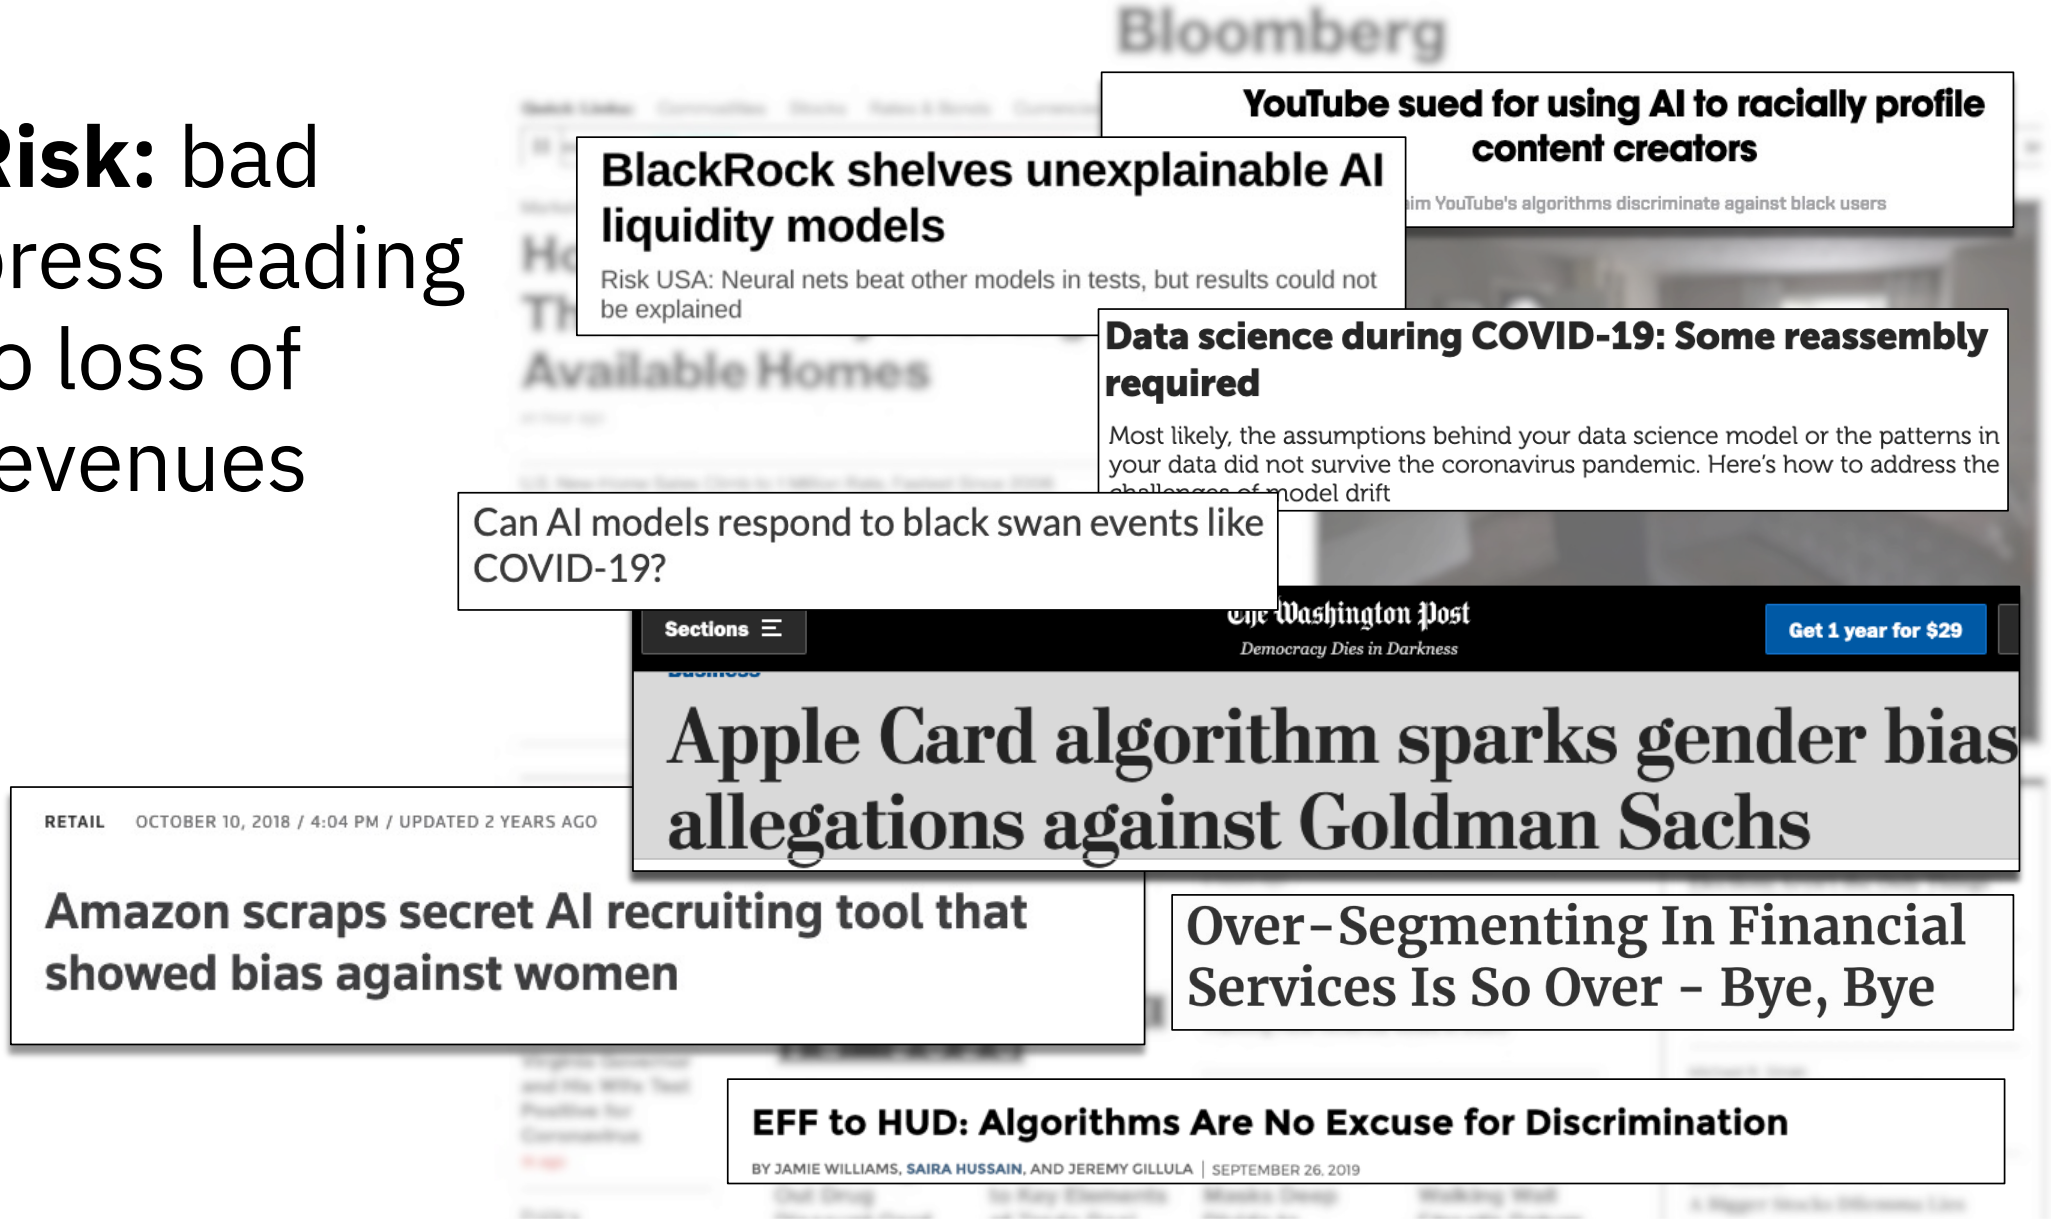

What does it take to trust a generative AI platform?

How was it trained?

- Garbage in -> garbage out
- An enterprise cannot use a foundation model trained with a Wikipedia crawl
- The training material needs to be huge and comprehensive but must also be curated

Can it detect & minimize bias & hallucination?

- How does the platform detect and correct bias?
- How can it prevent hallucination (providing random and untrue answers with absolute aplomb and convictions)?

Is it transparent?

- Open vs black-box
- How to **audit**, and explain the model and the answers it generates?
- Does the model track **drift and bias**? And how does it address them?

Does it support regulatory compliance?

- How do foundation models and their usage comply with privacy and government regulations?
- What are the guardrails?
- Who is responsible for an inadvertently exposed PII or a “wrong answer”?

Is it safe?

- Who has control over the model, input data, and output data?
- How to ensure that confidential information is not given out?
- How is it monitored?
- What safety features and guardrails are in place?

Can it be customized?

- Hybrid and multicloud?
- Can the model be fine-tuned with clients’ data?
- How can clients update, and extend the model to make it more suitable for their use cases?
- How to integrate with applications? What APIs are in place?

Put AI to work with watsonx

Scale and accelerate the impact of AI with trusted data

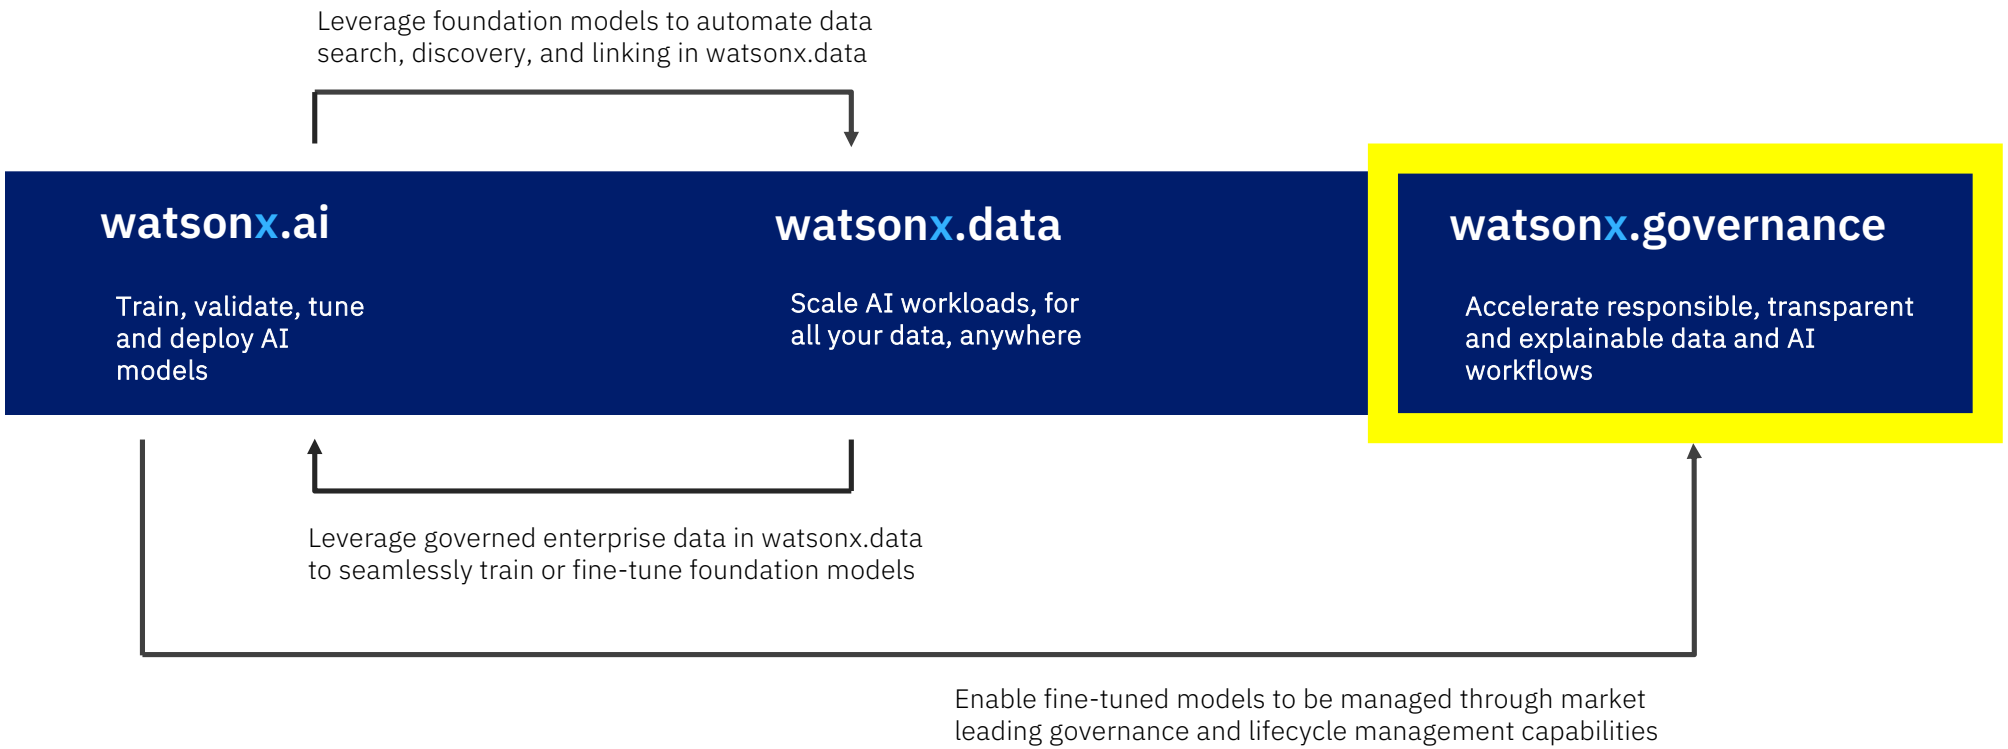

IBM watsonx.governance
↓
end-to-end automated AI lifecycle governance toolkit built to mitigate risk and improve compliance

watsonx.governance

Enable responsible, transparent and explainable AI workflows

An end-to-end AI lifecycle governance toolkit encompassing both data and AI governance to mitigate risk and improve compliance

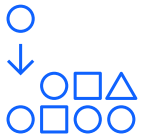

Govern across the AI lifecycle by automating and consolidating tools, applications and platforms.

Manage risk and protect reputation by automating workflows to better detect fairness, bias and drift.

Adhere to regulatory compliance by translating growing regulations into enforceable policies.

Comprehensive
Govern the end-to-end AI lifecycle with metadata capture at each stage

Open
Support governance of models built and deployed in 3rd party tools.

Automatic metadata
and data transformation/lineage capture though Python notebooks.

A Governed, Trusted AI Lifecycle

Workflows, dashboards for risk & incident management

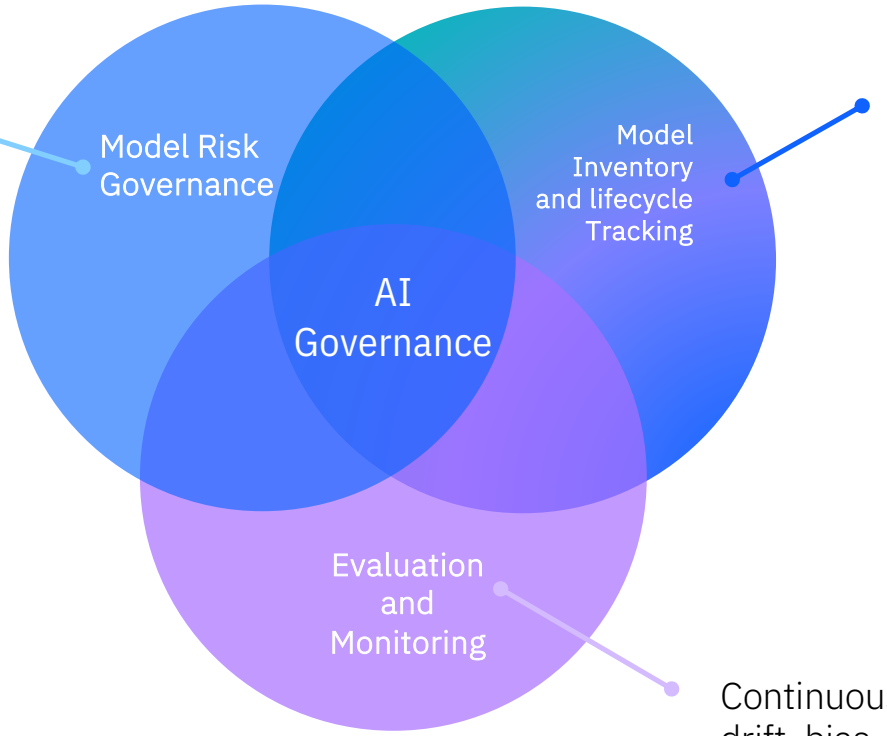

Capture model facts throughout the entire lifecycle

Continuous monitoring for accuracy, drift, bias, explainability, and performance

AI Governance based on watsonx.governance integrates and augments your existing development and deployment tools and processes

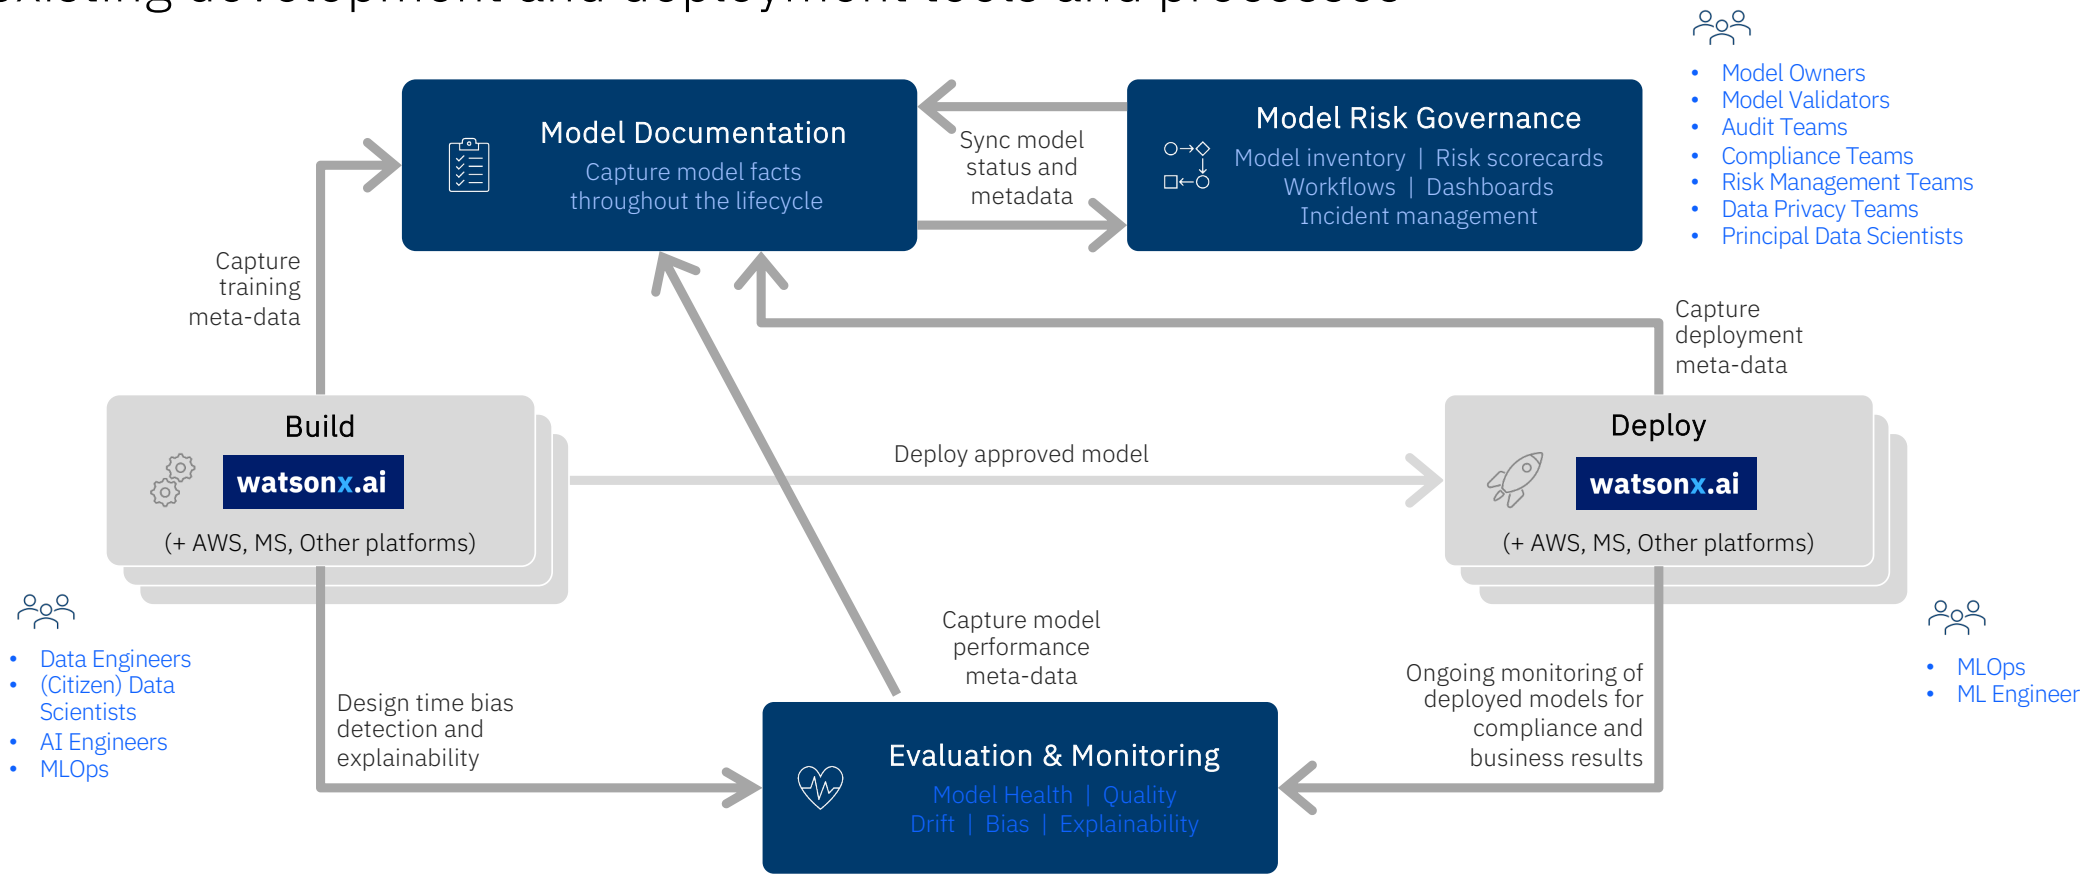

Model Inventory & Lifecycle Management

- Captures Model facts throughout the lifecycle.
- Provide a singular view of facts across the model lifecycle
- Facilitate subsequent enterprise validation, understand how the model will behave in different business situations
- Support audits, and requests for model facts from auditors, management, stakeholders, and customers

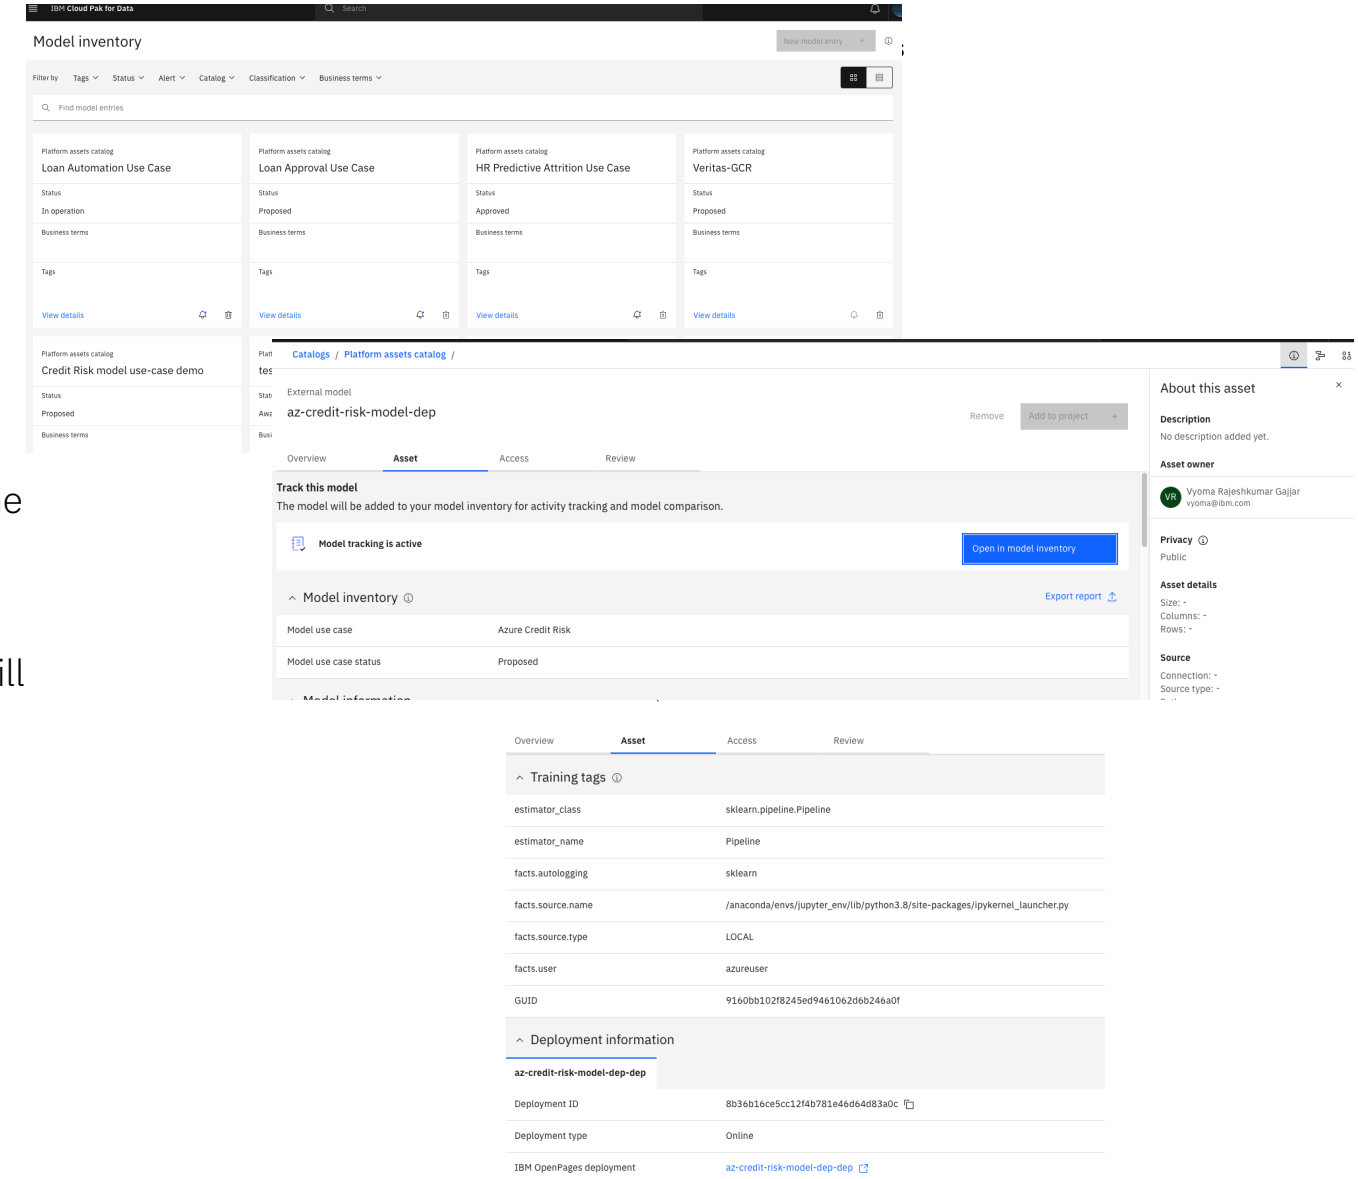

Evaluation and Monitoring

- Ongoing health monitoring of Models
- Trace and explain Model predictions
- Document metrics and track metric values over time
- Bias detection and mitigation
- Notification of issues when quality thresholds or business KPIs are violated

OpenScale

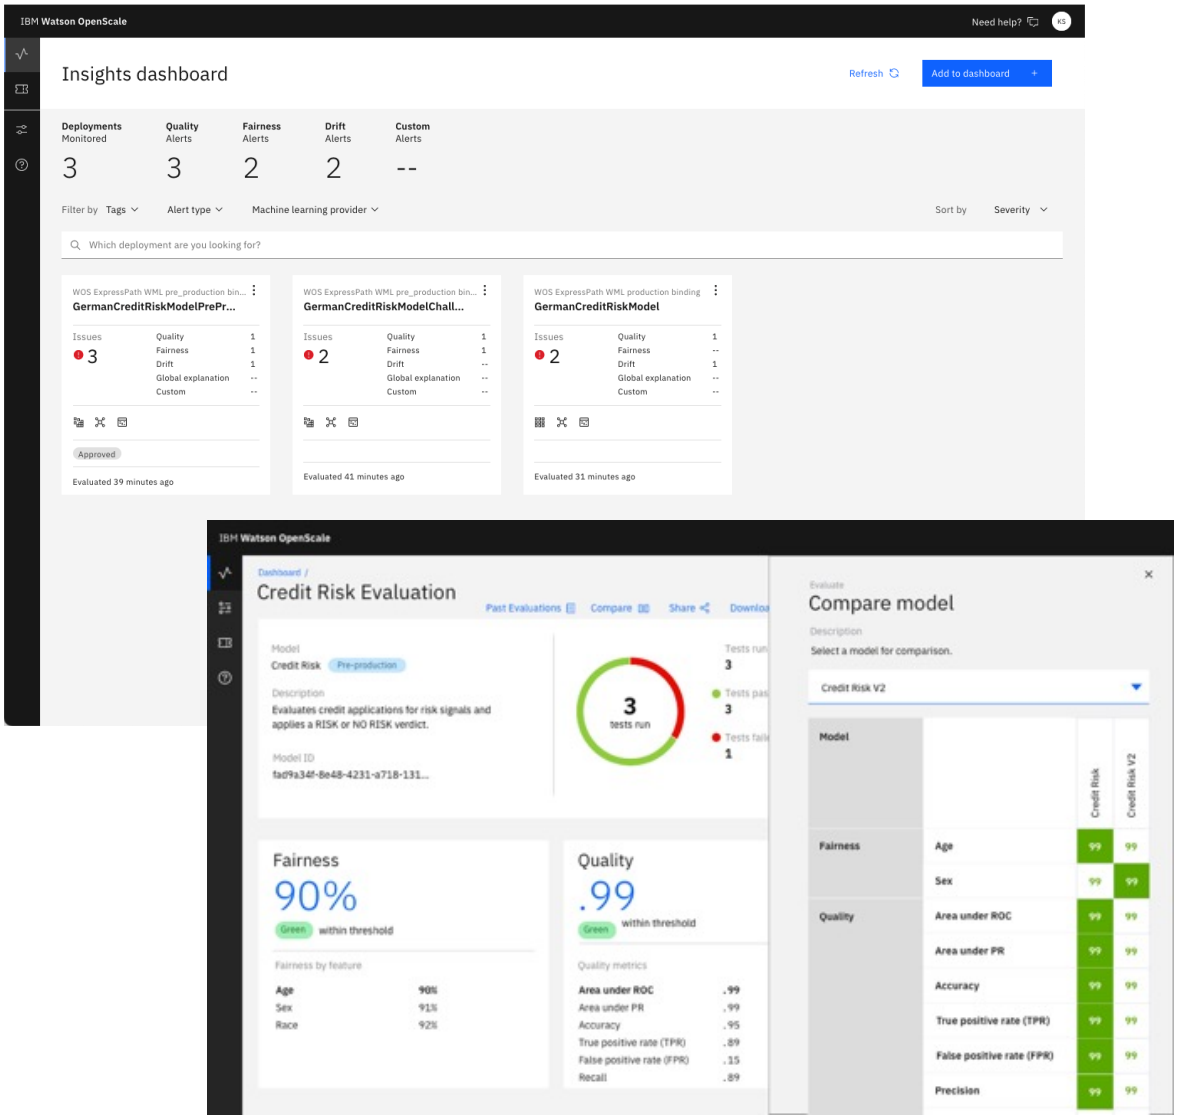

Model Risk Governance

1. Consolidated view of models from multiple platforms
2. View development status, model performance and alerts or emerging issues
3. Monitor and trigger workflows for model validation, retraining and performance issues
4. Track issues and incidents related to models in OpenPages included Issue Management Solution
5. Workflow to document and approve changes to models

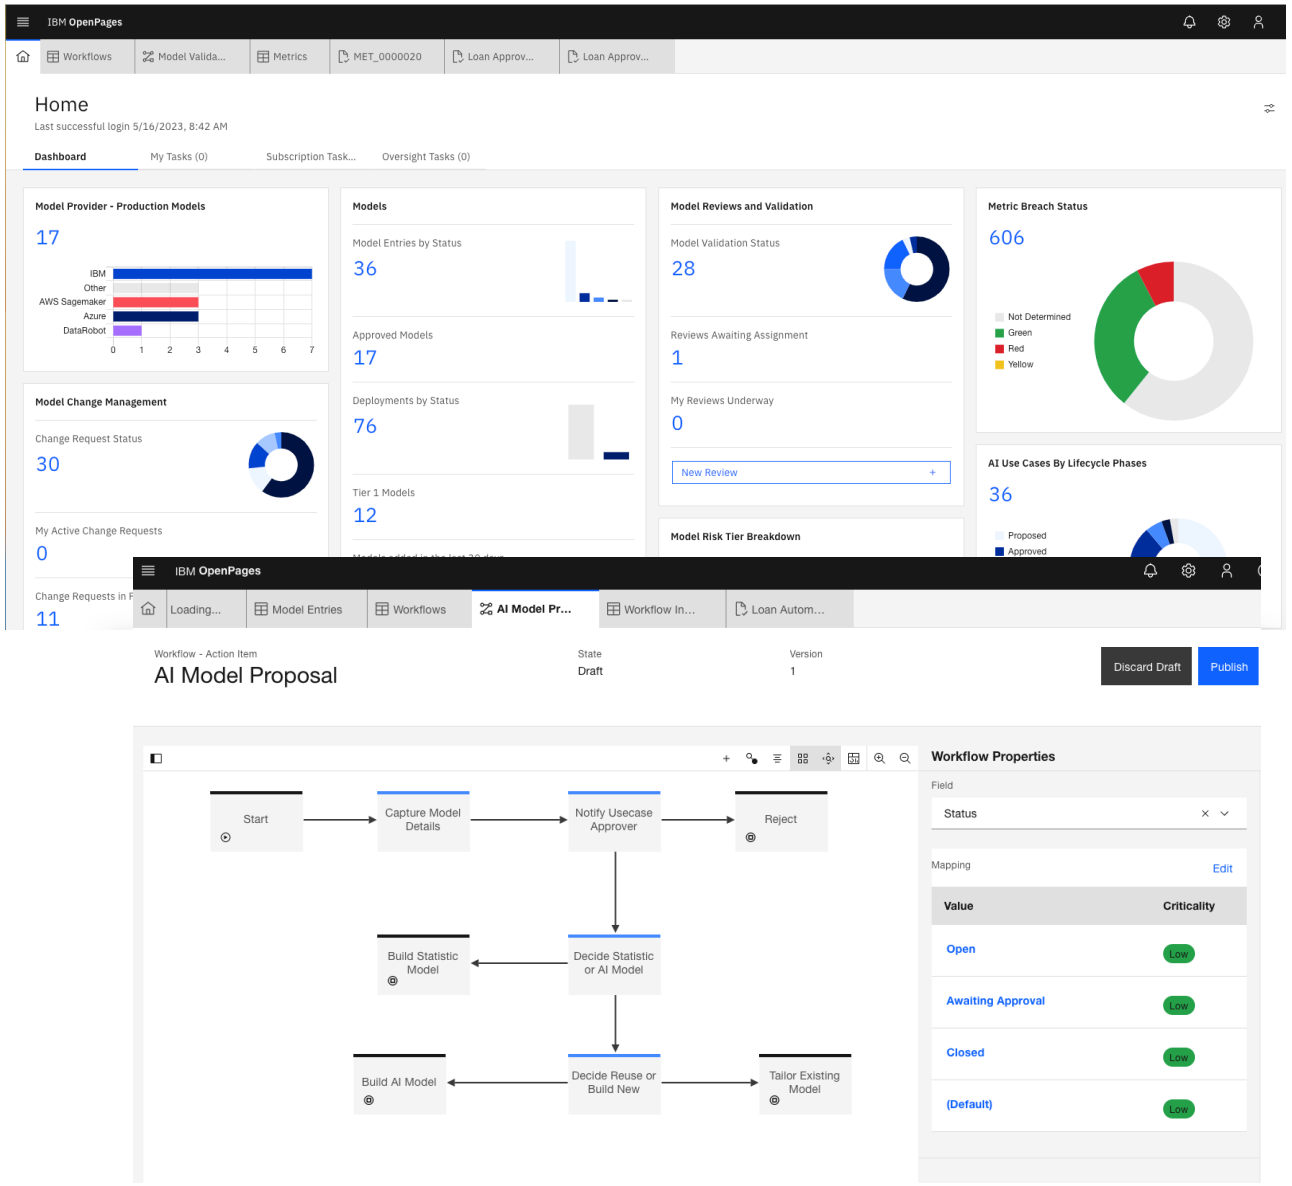

AI Lifecycle Governance – Process View

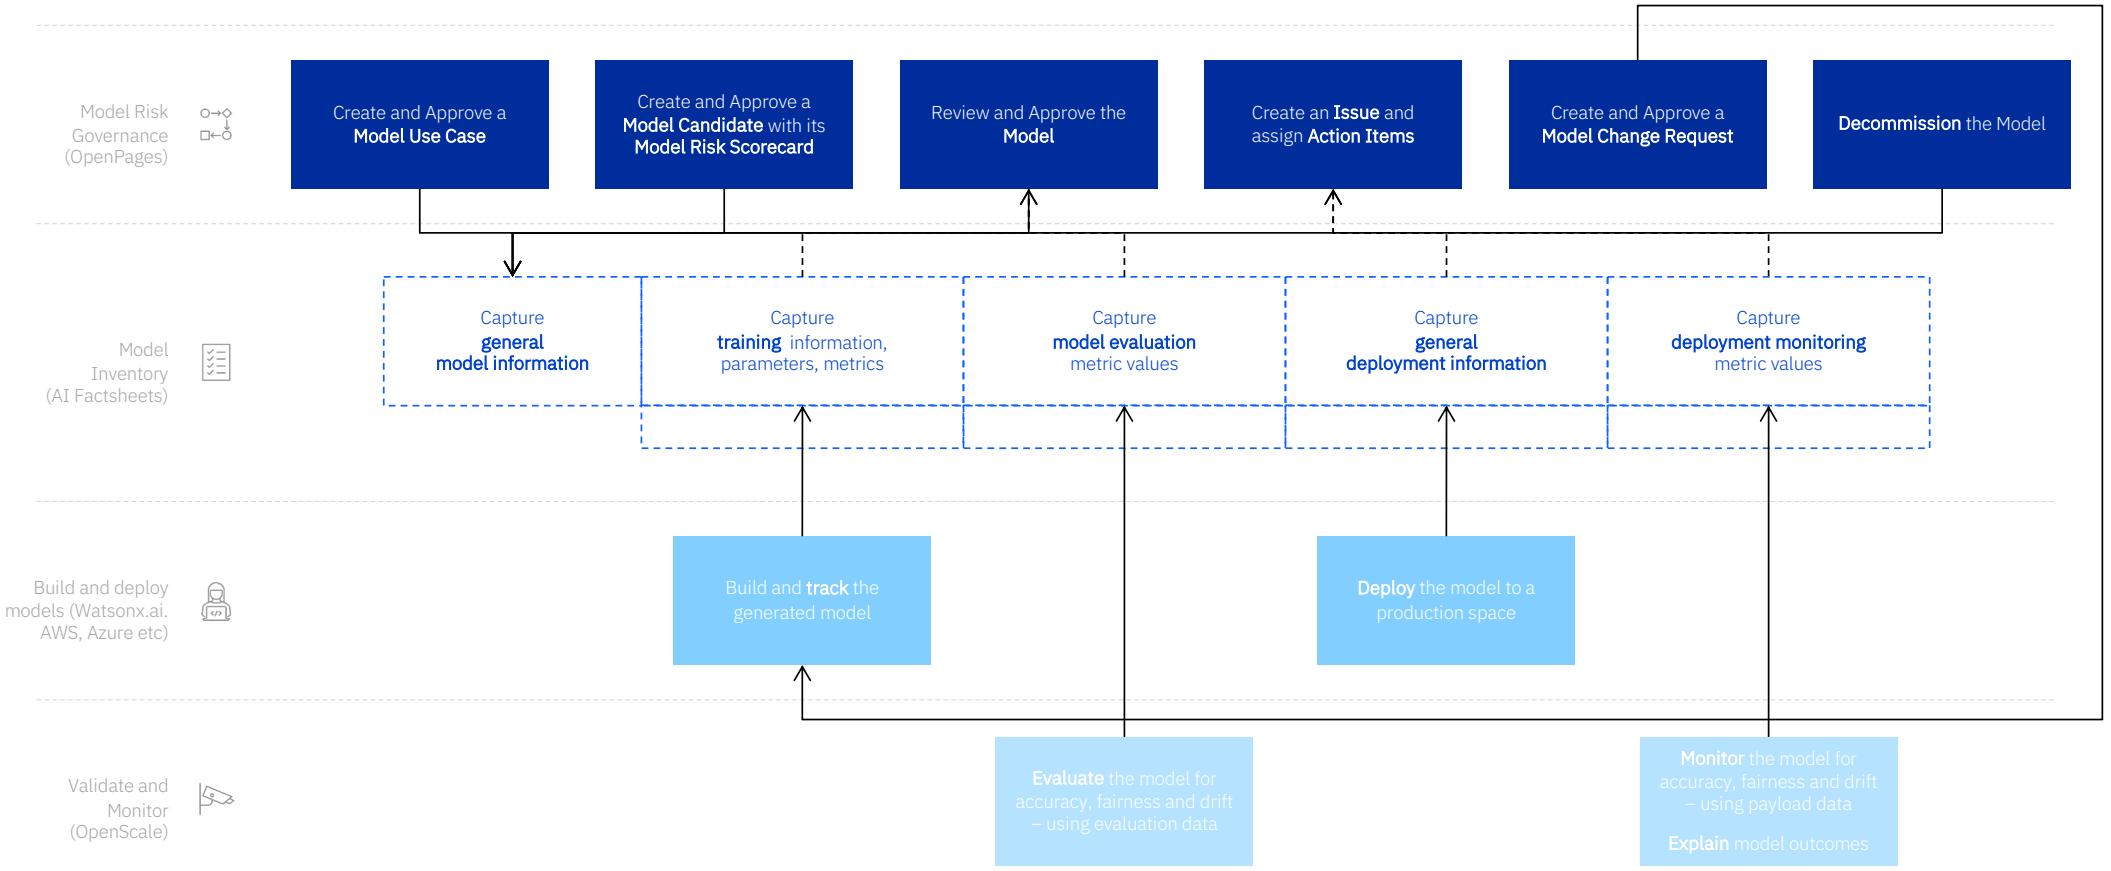

(boxes with solid fill colors are user actions, boxes with dashed outlines are what the integrated tooling does automatically)

Reference Architecture

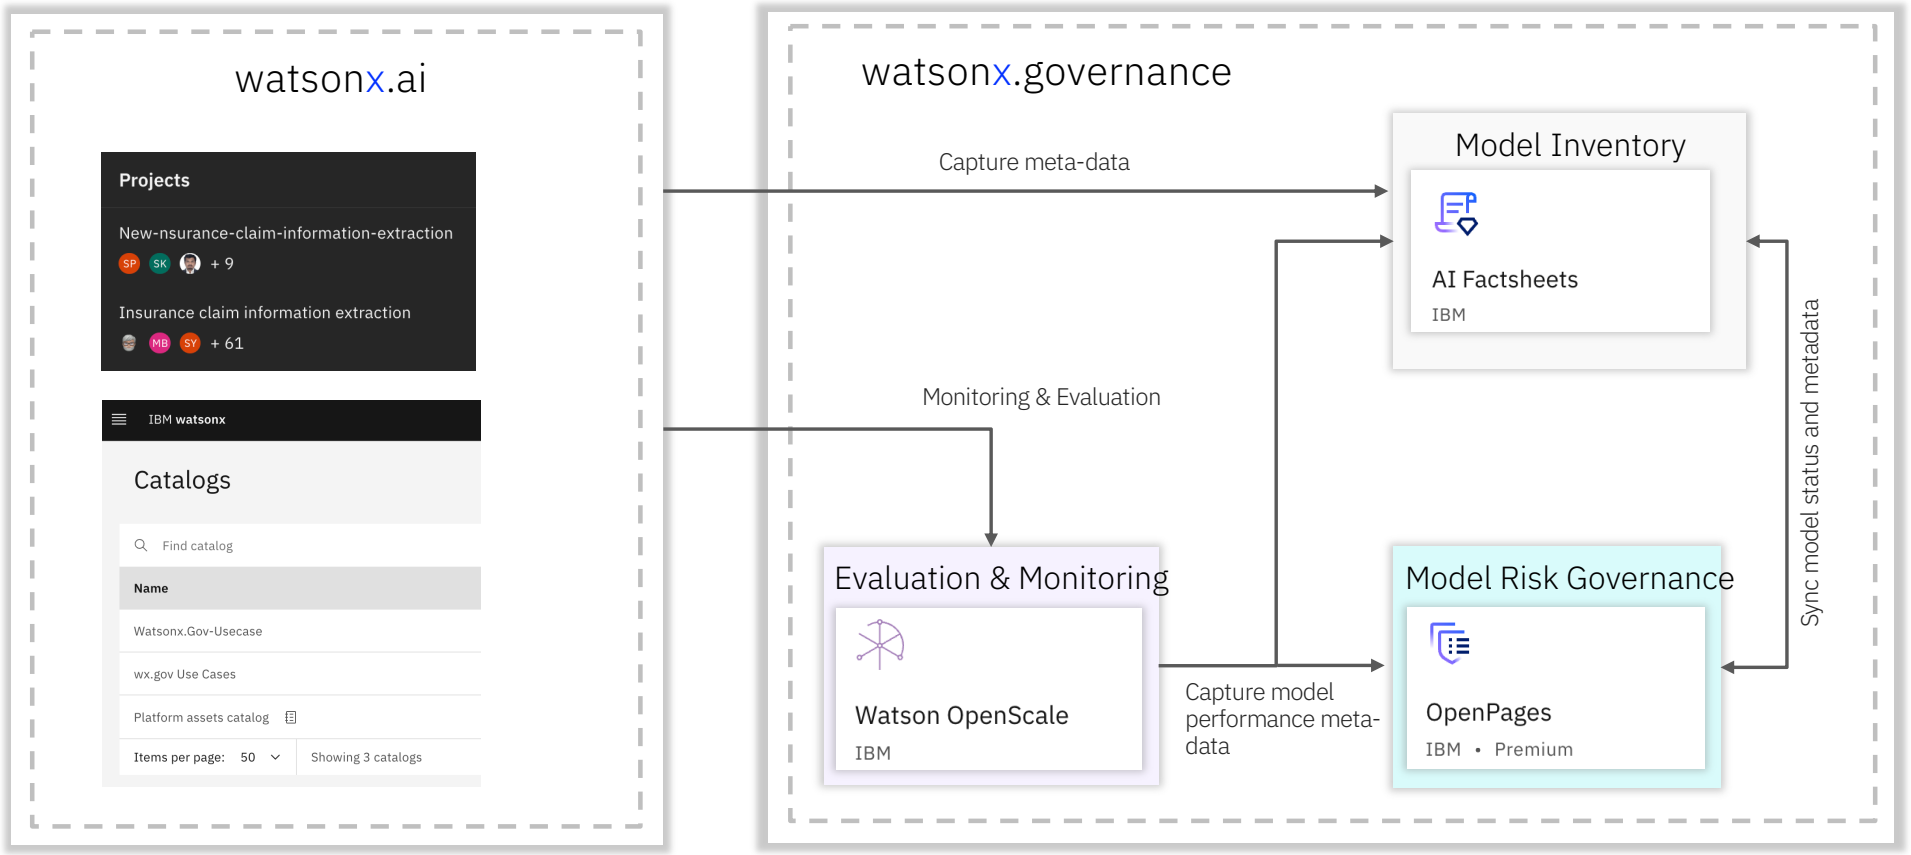

Stakeholder Pain Points

Head of Model Risk (Alex)

I can't keep track of the **thousands AI models and how we are tracking them and making sure model does not Risk on day 2.**

Enterprise Risk Director (Mary)

We have difficulty in effectively **managing AI models** introduced by evolving new business requirements

Head of Model Validation (Doe)

We have **fragmented and disparate practices** for managing models and metrics to measure their effectiveness across the organization

Prompt Engineer (Ann)

I need a single view to track all the crucial model parameters for multiple prompts on multiple models.

Business owner (Lisa)

I need the ability to see organizations 360 degree view of all models

Chief Compliance Officer (Jane)

I need the ability to define **multiple risk and compliance frameworks like EU AI Ethics act with harmonized controls around AI models**

Demo